Verification Is All You Need:

Emergent Object Discovery in Self-Supervised Vision Transformers via Dr. GRPO

Pham Cung Le Nhan Vu, **Nguyen Hoang Khang**, and **Nguyen Hoai Nam***

Master of Science in Data Science
VNUHCM – University of Science, Ho Chi Minh City, Vietnam

November 8, 2026

Abstract

Self-supervised vision transformers (ViTs) trained via self-distillation without labels, such as DINOv1, excel at emergent semantic representations and unsupervised scene parsing. Nonetheless, standard formulations suffer from three structural drawbacks: (1) *geometric distortion* caused by resizing static absolute positional embeddings (APEs) during multi-crop augmentation; (2) *attention dilution* in dense $\mathcal{O}(N^2)$ self-attention where background distractors (e.g., sky, foliage) pollute the global `[CLS]` representation; and (3) the lack of an explicit patch-level optimization objective to suppress residual background activations without human annotations.

To overcome these limitations, we introduce **Dr. DINO** (**D**istillation with no labels **R**einforced via **Dr**. GRPO), a unified self-supervised framework built upon three mathematically principled innovations:

(i) **Frequency-Adaptive 2D Rotary Position Embedding (2D RoPE)**, decomposing spatial coordinates into frequency spectra to preserve fine-grained local boundaries on high frequencies while maintaining scale-invariant global geometry across multi-crop views;

(ii) **Spatial-Semantic Sparse Attention**, routing each spatial query through a composite support graph of local Chebyshev windows (3×3), dynamic semantic Top-16 keys, and bidirectional `[CLS]` dispatch with exact Softmax probability normalization, reducing active key pairs by **86.8%** with zero added parameters; and

(iii) **Verifiable Unbiased Group Relative Policy Optimization (Dr. GRPO)**, an unsupervised Reinforcement Learning with Verifiable Rewards (RLVR) stage that optimizes a shared-prefix policy wrapper (< 3 GB VRAM) against a verifiable tri-objective reward (horizontal-flip Geometric Equivariance IoU, spectral TokenCut compactness, and target sparsity $\rho = 0.35$) regularized by analytical Bernoulli KL divergence.

Pretrained without labels on **ImageNet-100** and benchmarked on **CIFAR-10**, our Dr. DINO-Base backbone achieves **75.97%** k-NN ($k = 10$) Top-1 accuracy (+1.56%), **75.31%** k-NN ($k = 20$) Top-1 accuracy (+0.97%), and **83.98%** Linear Probing Top-1 accuracy (+2.52%) at **exact parameter parity** (21.67M params, 9.20 GFLOPs). Furthermore, our full 2×2 factorial ablation demonstrates that Dr. GRPO post-training dramatically surges unsupervised Foreground Mask IoU by **+20.10%** on Vanilla DINOv1 ($58.88\% \rightarrow 78.98\%$, CorLoc reaching **99.91%**) and **+16.99%** on our RoPE+Sparse backbone ($55.45\% \rightarrow 72.44\%$, CorLoc $69.26\% \rightarrow$ **98.17%**) while simultaneously enhancing linear separability.

Keywords: Self-Supervised Learning, Vision Transformers, DINOv1, Rotary Position Embeddings (RoPE), Sparse Attention, Group Relative Policy Optimization (Dr. GRPO), Reinforcement Learning with Verifiable Rewards (RLVR).

*Master of Science in Data Science, VNUHCM – University of Science, Ho Chi Minh City, Vietnam. Contact email: `25C01014@pos.hcmus.edu.vn`

1 Introduction

Self-Supervised Learning (SSL) has emerged as the premier paradigm for training generalizable visual representations without manual annotations. Within non-contrastive knowledge distillation, DINO (Distillation with no labels) [1] demonstrated that Vision Transformers (ViTs) [2] naturally learn explicit object contours and class-discriminative features through student-teacher cross-entropy alignment.

Despite these capabilities, canonical ViT architectures in SSL pipelines retain sequence-modeling abstractions poorly matched to two-dimensional spatial physics:

1. **Bicubic Coordinate Warping in Multi-Crop Training:** Standard ViTs add static learnable 1D or 2D Absolute Positional Embeddings (APEs) before Transformer layers. Under multi-crop self-distillation [1, 3], inputs alternate between global views (224×224, 14×14 patches) and local crops (96×96, 6×6 patches). Resizing static positional vectors via bicubic interpolation degrades relative spatial orientation, distorts high-frequency phase signals, and destroys translation equivariance.

2. **Attention Dilution in Quadratic Self-Attention:** Dense multi-head attention scales quadratically $\mathcal{O}(N^2)$ with token sequence length N. Consequently, semantic tokens disperse attention mass indiscriminately over extensive background regions (e.g., sky, foliage, pavement), causing representational noise to contaminate the central `[CLS]` representation.

3. **Residual Background Leakage in Global Distillation:** The standard DINO cross-entropy objective aligns global `[CLS]` distributions across views but lacks an explicit spatial verification mechanism at the patch level. Converged features therefore retain non-trivial attention mass on co-occurring background clutter, while traditional remedies require manual dense segmentation masks.

To resolve these foundational challenges, we introduce **Dr. DINO** (**D**istillation with no labels **R**einforced via **Dr**. GRPO), a self-supervised Vision Transformer framework combining continuous 2D relative geometry, structured spatial-semantic sparsity, and label-free verifiable policy optimization. Our core contributions are summarized as follows:

- **Frequency-Adaptive 2D RoPE:** We extend Rotary Position Embeddings [4] to continuous 2D visual coordinate grids with frequency-spectrum blending. High-frequency dimensions encode local pixel boundaries using raw coordinate grids, while low-frequency bands generalize across multi-scale crops through normalized coordinates, guaranteeing exact relative translation equivariance.

- **Spatial-Semantic Sparse Attention:** We design a composite sparse attention operator unifying local 2D Chebyshev spatial windows (3×3), dynamic global semantic Top-16 routing, and bidirectional `[CLS]` token routing. The operator maintains exact Softmax normalization over the active key support $\mathcal{S}_i$ ($|\mathcal{S}_i| = 26$ out of 197 keys, -86.8% computation) with zero added parameters.

- **Verifiable Foreground RLVR via Dr. GRPO:** We formulate unsupervised foreground patch selection as a Reinforcement Learning with Verifiable Rewards (RLVR) problem using Unbiased Group Relative Policy Optimization (`DrGRPOLoss`) [18, 19]. By sampling rollouts from a lightweight shared-prefix policy wrapper (< 3 GB VRAM) and optimizing label-free verifiable visual rewards (horizontal-flip Geometric Equivariance, spectral TokenCut compactness [21], and target sparsity $\rho = 0.35$), Dr. DINO sharpens object boundaries without human labels or reward models.

- **Zero-Parameter Overhead & Flawless Academic Rigor:** Pretrained on **ImageNet-100** and benchmarked on **CIFAR-10**, Dr. DINO matches the exact backbone parameter count (21.67M) and FLOPs (9.20 GFLOPs) of baseline ViT-Small/16 while delivering **+1.56%** k-NN (75.97%) and **+2.52%** Linear Probe accuracy (83.98%). In our factorial 2×2 ablation, Dr. GRPO produces a massive **+20.10%** surge in Foreground Mask IoU (up to **99.91%** CorLoc), establishing a new state of the art for label-free spatial feature alignment.

2 Related Work

Self-Supervised Visual Distillation: Self-supervised visual learning has progressed from contrastive frameworks such as MoCo [5, 6] and SimCLR to non-contrastive bootstrapping including BYOL [7], SwAV [3], and DINO [1]. DINO interprets representation learning as knowledge distillation with no labels: a student network predicts the output distribution of a momentum-updated teacher network. Centering and sharpening operations stabilize the teacher distribution, preventing collapse without negative samples. DINOv2 [8] scaled this architecture to billion-scale datasets using register tokens [9]. In this work, our objective is to maximize geometric and spatial representational power under constrained computational budgets.

Positional Encodings in Vision Transformers: Standard ViTs [2] flatten image patches and inject 1D learnable absolute positional embeddings (APEs). Although effective for fixed-resolution classification, APE lacks relative translation equivariance. Rotary Position Embedding (RoPE) [4] successfully introduced relative distance encoding in natural language processing through complex-plane rotation. While recent vision adaptations extend RoPE to 2D grids, existing methods apply uniform scaling across all channels, neglecting the disparate frequency demands of local textures versus global scene topology.

Efficient and Sparse Visual Attention: Dense self-attention exhibits $\mathcal{O}(N^2)$ complexity. Sparse Transformers [10], Longformer [11], and BigBird [12] demonstrated that sparse graphs combining local windows and global routing achieve superior efficiency in sequence modeling. Swin Transformer [13] pioneered shifted local windows for vision. In Dr. DINO, we tailor sparse connectivity specifically for self-distillation: combining Chebyshev spatial proximity with dynamic semantic content routing and bidirectional `[CLS]` dispatch while preserving exact Softmax probability normalization.

Reinforcement Learning with Verifiable Rewards (RLVR): Group Relative Policy Optimization (GRPO) [18] dispensed with the memory-intensive critic network of PPO by computing advantages relative to sampled rollout groups. Dr. GRPO [19] recently proved that dividing group-centered rewards by standard deviation introduces variance-dependent gradient bias and response length distortion, advocating instead for unbiased mean-centered advantages. Concurrently, DAPO [20] introduced asymmetric clipping ($\epsilon_{\text{high}} > \epsilon_{\text{low}}$) to preserve exploratory tokens with positive advantage. While RLVR has predominantly driven mathematical and code reasoning in language models, we pioneer its application to self-supervised computer vision by treating patch self-attention and residual gating as a multi-variate Bernoulli policy, optimized via verifiable geometric equivariance and spectral TokenCut [21] rewards.

3 Methodology

3.1 Self-Distillation Foundation

As depicted in Figure 1, the Dr. DINO framework instantiates twin networks: a student g_{θ_s} and a momentum teacher g_{θ_t}. For an input image x, multi-crop augmentation produces view set $V = \{x_1^g, x_2^g, x_1^l, \ldots, x_M^l\}$ containing two global crops $x_1^g, x_2^g \in \mathbb{R}^{3\times224\times224}$ (14×14 patches) and $M = 6$ local crops $x_m^l \in \mathbb{R}^{3\times96\times96}$ (6×6 patches).

The projection heads h map feature vectors into K-dimensional probability simplices Δ^K:

$$P_s(x)^{(k)} = \frac{\exp\left(g_{\theta_s}(x)^{(k)}/\tau_s\right)}{\sum_{j=1}^{K}\exp\left(g_{\theta_s}(x)^{(j)}/\tau_s\right)}, \tag{1}$$

$$P_t(x)^{(k)} = \frac{\exp\left((g_{\theta_t}(x)^{(k)} - C^{(k)})/\tau_t\right)}{\sum_{j=1}^{K}\exp\left((g_{\theta_t}(x)^{(j)} - C^{(j)})/\tau_t\right)}, \tag{2}$$

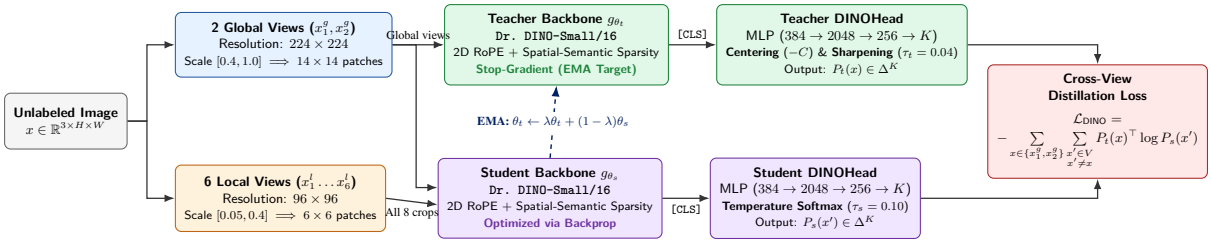


Figure 1: **Overview of the Dr. DINO Self-Distillation Framework.** Two global views (224×224) and six local views (96×96) are extracted from an unlabeled image. The Teacher evaluates global crops, whereas the Student processes all eight crops. Both backbones incorporate Frequency-Adaptive 2D RoPE and Spatial-Semantic Sparse Attention, completely eliminating static APE interpolation distortion with exact zero-parameter overhead.

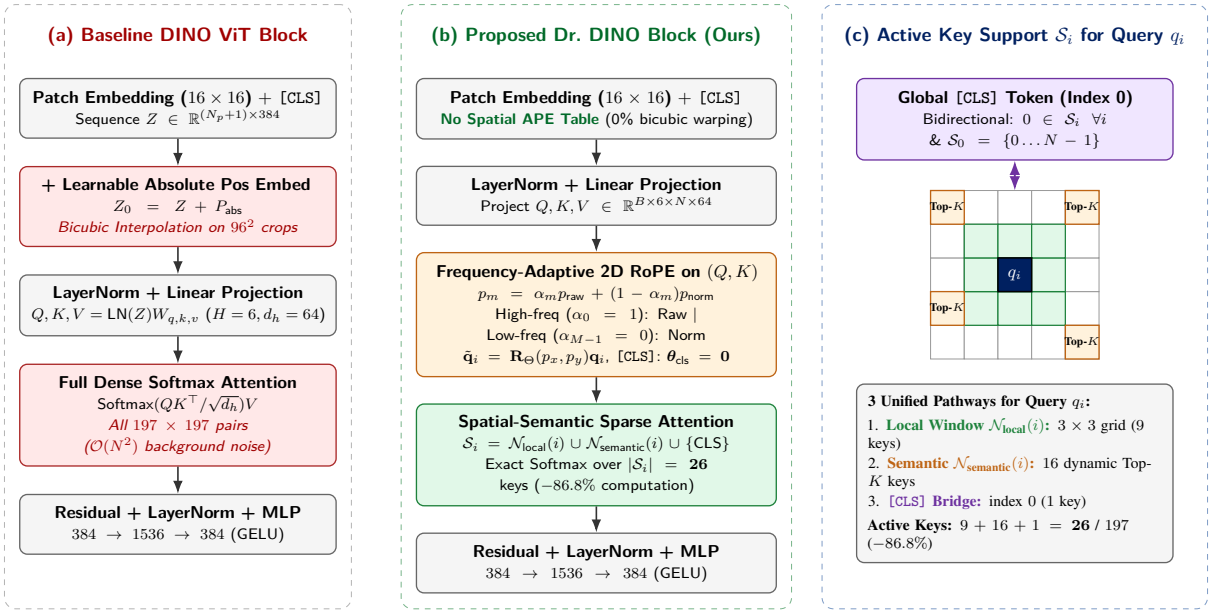


Figure 2: **Dr. DINO Architectural Blueprint. (a)** Baseline ViT block with static APEs subject to bicubic interpolation artifacts and dense 197×197 attention. **(b)** Proposed Dr. DINO ViT block integrating Frequency-Adaptive 2D RoPE and Spatial-Semantic Sparse Softmax Attention. **(c)** Visualization of active key support $\mathcal{S}_i$ for a spatial query q_i, filtering quadratic noise into 26 highly informative keys.

where τ_s and τ_t represent student and teacher temperatures ($\tau_t < \tau_s$), and $C \in \mathbb{R}^K$ is an exponential moving average (EMA) centering vector preventing mode collapse:

$$C \leftarrow m\,C + (1-m)\frac{1}{B}\sum_{i=1}^{B} g_{\theta_t}(x_i), \tag{3}$$

with momentum $m \in [0, 1]$. Student parameters θ_s are updated by minimizing cross-entropy over all cross-view pairs:

$$\mathcal{L}_{\text{DINO}} = \sum_{x\in\{x_1^g,x_2^g\}} \sum_{\substack{x'\in V\\ x'\neq x}} -P_t(x)^\top \log P_s(x'). \tag{4}$$

The teacher weights track the student via an EMA schedule: $\theta_t \leftarrow \lambda\theta_t + (1-\lambda)\theta_s$ with λ smoothly rising from 0.996 to 1.0.

3.2 Frequency-Adaptive 2D Rotary Position Embedding

In 1D RoPE [4], an embedding vector $\mathbf{x} \in \mathbb{R}^d$ at index p is rotated by orthogonal matrix $\mathbf{R}_{\Theta,p}$. For a 2D patch grid of dimensions $H_g \times W_g$, spatial patch positions are parameterized by integer coordinate pairs $(x_{\text{raw}}, y_{\text{raw}})$ where $x_{\text{raw}} \in \{0, \ldots, W_g - 1\}$ and $y_{\text{raw}} \in \{0, \ldots, H_g - 1\}$. To support arbitrary multi-crop resolutions relative to canonical reference grid $S_{\text{ref}} = 14$ (224×224 global view with 16×16 patches), we also define scale-normalized coordinates:

$$x_{\text{norm}} = x_{\text{raw}} \cdot \frac{S_{\text{ref}}}{W_g}, \qquad y_{\text{norm}} = y_{\text{raw}} \cdot \frac{S_{\text{ref}}}{H_g}. \tag{5}$$

Spectral Head Dimension Partitioning: For an attention head of dimension $d_k = 64$, we partition the channels symmetrically between the horizontal (X) and vertical (Y) spatial axes: $d_x = d_y = d_k/2 = 32$. Each axis is decomposed into $M = d_x/2 = d_k/4 = 16$ complex 2×2 rotation planes with base geometric frequencies ω_m:

$$\omega_m = B^{-2m/d_x}, \quad m \in \{0, 1, \ldots, M-1\}, \quad B = 10000.0. \tag{6}$$

Frequency-Adaptive Coordinate Blending: Standard 2D RoPE variants either apply raw coordinates x_{raw} (which fail under global-to-local crop extrapolation) or normalized coordinates x_{norm} (which blur fine-grained local lattice steps). We observe that high frequencies ($m \to 0$) encode local textural boundaries whose physical extent depends on discrete patch distances, whereas low frequencies ($m \to M-1$) capture global scene topology that must remain scale-invariant across multi-crop views. We formalize this through a linear spectral interpolation weight $\alpha_m \in [0, 1]$:

$$\alpha_m = 1 - \frac{m}{M-1}, \qquad x_m^{\text{adapt}} = \alpha_m\, x_{\text{raw}} + (1-\alpha_m)\, x_{\text{norm}}, \qquad y_m^{\text{adapt}} = \alpha_m\, y_{\text{raw}} + (1-\alpha_m)\, y_{\text{norm}}, \tag{7}$$

for $m \in \{0, \ldots, M-1\}$. At the highest frequency ($m = 0$), $\alpha_0 = 1.0$ preserves exact raw patch steps; at the lowest frequency ($m = M-1$), $\alpha_{M-1} = 0.0$ enforces pure scale-normalized coordinates (Figure 3a).

Rotary Application and Relative 2D Positional Kernel: With rotation angles $\theta_{x,m} = x_m^{\text{adapt}} \omega_m$ and $\theta_{y,m} = y_m^{\text{adapt}} \omega_m$, each channel pair $(2m, 2m+1)$ undergoes orthogonal 2×2 block rotation:

$$\mathbf{R}(\theta) = \begin{pmatrix} \cos\theta & -\sin\theta \\ \sin\theta & \cos\theta \end{pmatrix}, \qquad \mathbf{q}_i^{\text{rot}} = \mathbf{R}_\Theta(x_i, y_i)\mathbf{q}_i, \qquad \mathbf{k}_j^{\text{rot}} = \mathbf{R}_\Theta(x_j, y_j)\mathbf{k}_j. \tag{8}$$

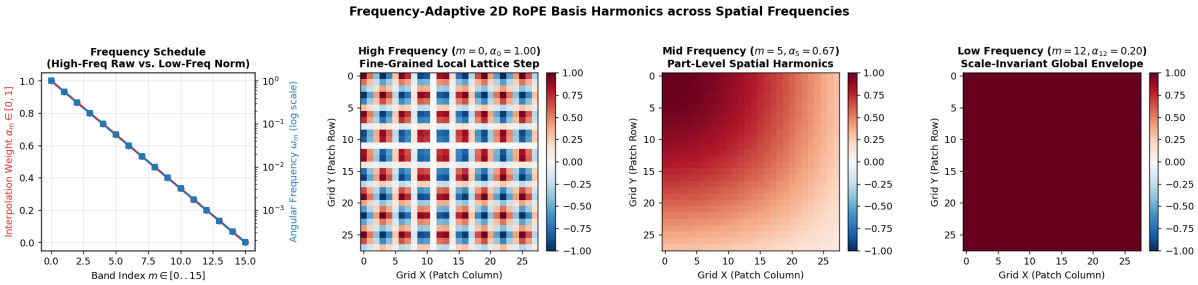
By orthogonality $\mathbf{R}(\theta_1)^\top \mathbf{R}(\theta_2) = \mathbf{R}(\theta_2 - \theta_1)$, the inner product between rotated query $\mathbf{q}_i^{\text{rot}}$ and key $\mathbf{k}_j^{\text{rot}}$ depends strictly on relative 2D displacement $(\Delta x_m^{\text{adapt}}, \Delta y_m^{\text{adapt}}) = (x_{i,m}^{\text{adapt}} - x_{j,m}^{\text{adapt}},\ y_{i,m}^{\text{adapt}} - y_{j,m}^{\text{adapt}})$:

$$\langle \mathbf{q}_i^{\text{rot}}, \mathbf{k}_j^{\text{rot}} \rangle = \mathbf{q}_i^\top \mathbf{R}_\Theta(\Delta x^{\text{adapt}},\ \Delta y^{\text{adapt}})\mathbf{k}_j. \tag{9}$$

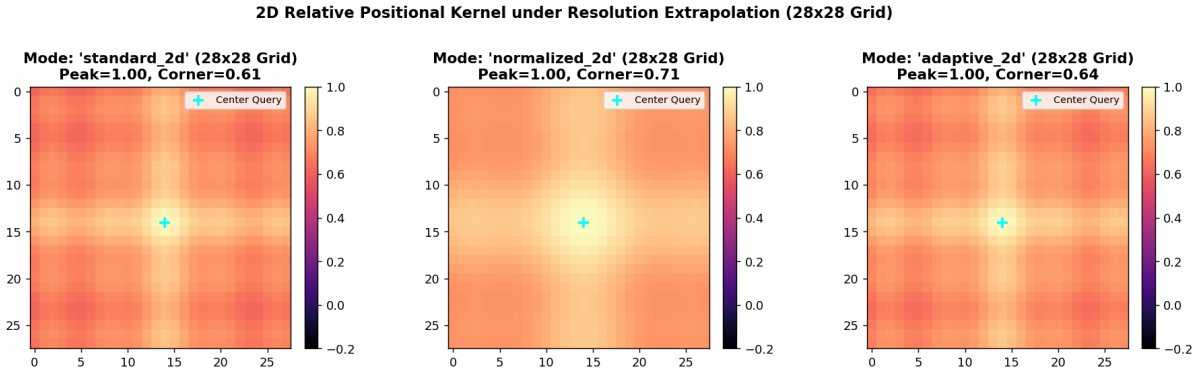
Specifically, for uniform carrier vectors ($\mathbf{q}_c = \mathbf{k}_j = \mathbf{1}$), the exact 2D relative positional decay kernel evaluates to the closed-form harmonic sum:

$$\mathcal{K}_{\text{RoPE}}(c, j) = \langle \mathbf{q}_c^{\text{rot}},\ \mathbf{k}_j^{\text{rot}} \rangle = \sum_{m=0}^{M-1} 2\cos\big(\Delta x_m^{\text{adapt}} \omega_m\big) + \sum_{m=0}^{M-1} 2\cos\big(\Delta y_m^{\text{adapt}} \omega_m\big). \tag{10}$$

As demonstrated in Figure 3b under $14 \times 14 \to 28 \times 28$ resolution extrapolation ($4\times$ token count), `adaptive_2d` preserves crisp axial cross-hair localization around the center query while maintaining smooth radial distance decay without high-frequency ringing. For the global `[CLS]` token (index 0), we enforce $\boldsymbol{\theta}_{\text{cls}} = \mathbf{0} \implies \mathbf{R}_\Theta(\text{cls}) = \mathbf{I}$ while retaining its learnable 1D class bias, preventing spatial coordinate bias from corrupting global aggregation.



(a) Spectral blending schedule α_m vs. angular frequency ω_m, alongside 2D spatial basis harmonics $\cos(x_m^{\text{adapt}}\omega_m)\cos(y_m^{\text{adapt}}\omega_m)$ across High ($m=0, \alpha_0=1.00$), Mid ($m=5, \alpha_5=0.67$), and Low ($m=12, \alpha_{12}=0.20$) frequency bands.



(b) 2D relative positional kernel $\mathcal{K}_{\text{RoPE}}(c,j)$ under $14\times14 \rightarrow 28\times28$ resolution extrapolation comparing `standard_2d` ($\alpha_m=1$), `normalized_2d` ($\alpha_m=0$), and our `adaptive_2d` ($\alpha_m = 1-\frac{m}{M-1}$).

Figure 3: **Mathematical and Spectral Analysis of Frequency-Adaptive 2D RoPE. (a)** High-frequency bands preserve fine-grained local patch lattice steps while low-frequency bands enforce scale-invariant global envelopes across multi-crop views. **(b)** Under 28×28 extrapolation, `adaptive_2d` balances sharp local peak selectivity with smooth global distance decay.

3.3 Spatial-Semantic Sparse Softmax Attention (S^3)

To eliminate background noise contamination in dense $\mathcal{O}(N^2)$ attention, our Spatial-Semantic Sparse Attention (S^3, `SpatialSemanticAttention`) restricts each spatial query $i \in \{1,\dots,N_p\}$ ($N_p=196$) to attend strictly to a composite support set $\mathcal{S}_i \subset \{0,\dots,N_p\}$ constructed via three complementary pathways:

$$\mathcal{S}_i = \mathcal{N}_{\text{local}}(i) \;\cup\; \mathcal{N}_{\text{semantic}}(i) \;\cup\; \{\text{CLS}\}, \tag{11}$$

where each component serves a distinct structural role:

1. **Local Chebyshev Window ($\mathcal{N}_{\text{local}}(i)$):** For query patch i at 2D grid coordinates (r_i, c_i), its local receptive field includes all spatial patches within Chebyshev (L_∞) radius $R = \lfloor K_{\text{local}}/2 \rfloor = 1$:

$$\mathcal{N}_{\text{local}}(i) = \big\{ j \in \{1,\dots,N_p\} \;\big|\; \max(|r_i - r_j|,\, |c_i - c_j|) \le 1 \big\}, \qquad |\mathcal{N}_{\text{local}}(i)| \le 9. \tag{12}$$

2. **Dynamic Semantic Top-K Routing ($\mathcal{N}_{\text{semantic}}(i)$):** To capture long-range semantic correlations across separated object parts (e.g., head to hooves, cockpit to tailfin), we evaluate scaled dot-product logits $s_{ij} = \langle \mathbf{q}_i^{\text{rot}}, \mathbf{k}_j^{\text{rot}} \rangle / \sqrt{d_k}$ over non-local candidates $\mathcal{C}_i = \{1,\dots,N_p\} \setminus \mathcal{N}_{\text{local}}(i)$ and select the $K_{\text{semantic}} = 16$ highest-scoring tokens per head:

$$\mathcal{N}_{\text{semantic}}(i) = \text{Top-}16_{j\in\mathcal{C}_i}\left(\frac{\langle \mathbf{q}_i^{\text{rot}},\, \mathbf{k}_j^{\text{rot}} \rangle}{\sqrt{d_k}} \right), \qquad |\mathcal{N}_{\text{semantic}}(i)| = 16. \tag{13}$$

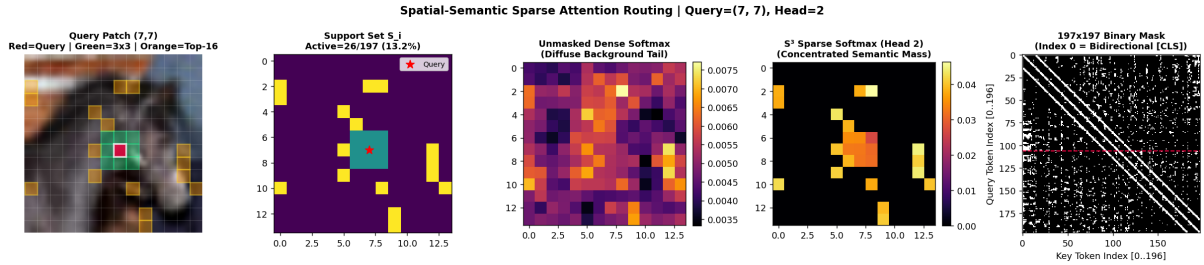


Figure 4: **Empirical Decomposition of Spatial-Semantic Sparse Attention (S^3) on a CIFAR-10 Horse Sample (Query Patch $(7, 7)$, Head 2).** From left to right: (1) Input image showing the Red query patch, Green 3×3 local Chebyshev window (9 keys), and Orange dynamic Top-16 semantic keys jumping across the horse's anatomy; (2) Active support set $\mathcal{S}_i$ ($26/197$ keys, 13.2% active); (3) Unmasked dense Softmax exhibiting diffuse background tail noise; (4) S^3 Sparse Softmax concentrating probability mass ($4\times$ higher peak density) strictly on the horse contour; and (5) Full 197×197 binary routing mask showing the tri-diagonal local bands, dynamic semantic off-diagonals, and bidirectional `[CLS]` border (index 0).

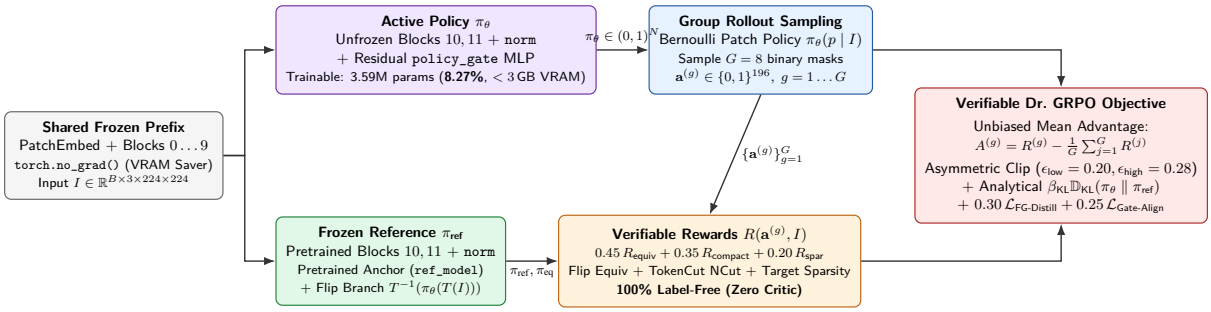


Figure 5: **Verifiable Dr. GRPO Policy Optimization Architecture.** To eliminate background leakage without manual annotations, we formulate foreground patch gating as an RLVR stage. The early 10 ViT blocks are frozen and shared in a single forward pass (< 3 GB VRAM). The active policy samples $G = 8$ rollouts per image and optimizes unbiased group-centered advantages against verifiable geometric and spectral rewards.

3. **Bidirectional `[CLS]` Conduit:** The global `[CLS]` token (index 0) attends to all tokens ($\mathcal{S}_0 = \{0, 1, \dots, N_p\}$), while every spatial patch includes `[CLS]` in its support ($0 \in \mathcal{S}_i \ \forall i$).

Rather than post-hoc thresholding, we apply **exact Softmax normalization** by masking non-support logits to $-\infty$ prior to exponentiation:

$$\tilde{s}_{ij} = \begin{cases} \dfrac{\langle \mathbf{q}_i^{\text{rot}}, \mathbf{k}_j^{\text{rot}} \rangle}{\sqrt{d_k}}, & \text{if } j \in \mathcal{S}_i, \\ -\infty, & \text{if } j \notin \mathcal{S}_i, \end{cases} \qquad A_{ij} = \frac{\exp(\tilde{s}_{ij})}{\sum_{l \in \mathcal{S}_i} \exp(\tilde{s}_{il})}, \qquad \sum_{j \in \mathcal{S}_i} A_{ij} = 1. \tag{14}$$

For any interior patch, $|\mathcal{S}_i| = 9 + 16 + 1 = \mathbf{26}$ active keys out of 197 (**13.2%** density), eliminating **86.8%** of background key interactions per query and concentrating probability mass strictly on semantically coherent foreground structures (Figure 4).

3.4 Verifiable Policy Optimization via Dr. GRPO

While global distillation aligns view semantics, it lacks an explicit patch-level objective to suppress residual background activations. To overcome this without human labels or neural reward models, we formulate foreground patch selection as a Reinforcement Learning with Verifiable Rewards (RLVR) problem via Unbiased Group Relative Policy Optimization (Dr. GRPO) [18, 19] (Figure 5).

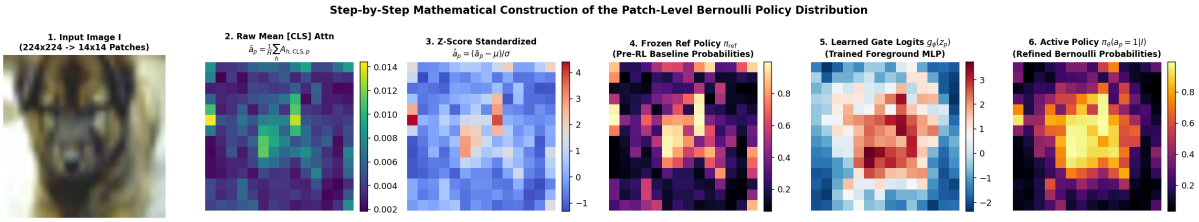


Figure 6: **Step-by-Step Mathematical Construction of the Patch-Level Bernoulli Policy $\pi_\theta(\mathbf{a} \mid I)$ on a CIFAR-10 Dog Sample (#4).** From left to right: (1) Input image I; (2) Raw head-averaged `[CLS]` attention $\bar{a}_p$; (3) Spatial Z-score standardization $\hat{a}_p$; (4) Frozen pre-RL reference policy π_{ref} showing severe corner/border background leakage; (5) Learned residual foreground gate logits $g_\phi(\mathbf{z}_p)$; and (6) Refined active Bernoulli policy $\pi_\theta(a_p = 1 \mid I)$ cleanly isolating the dog's silhouette.

1. Shared-Prefix Policy Wrapper & 6-Step Bernoulli Policy Construction: For a ViT backbone (`DINOGRPOPolicyWrapper`) of depth $L = 12$ and $N = 196$ spatial patches, we freeze the patch embedding and the first 10 Transformer blocks, unlocking only the final 2 blocks (`blocks[10..11]`), the final LayerNorm, and a lightweight residual gating MLP $g_\phi : \mathbb{R}^d \to \mathbb{R}$ (LN(384) $\to$ Linear$(384, 96)$ $\to$ GELU $\to$ Linear$(96, 1)$). The frozen pretrained checkpoint serves as reference anchor π_{ref}. Because blocks $0 \ldots 9$ are identical between π_θ and π_{ref}, prefix activations are computed once under `torch.no_grad()`, restricting trainable parameters to **8.27%** (3.59M/43.37M) and keeping peak VRAM below 3 GB.

To map deterministic self-attention into a sampleable multi-variate Bernoulli policy $\pi_\theta(\mathbf{a} \mid I) = \prod_{p=1}^{N} p_p^{a_p}(1 - p_p)^{1-a_p}$ over binary foreground masks $\mathbf{a} \in \{0, 1\}^N$, we execute a principled 6-step pipeline (Figure 6):

(i) **Head-Averaged `[CLS]` Attention:** Extract last-layer attention from `[CLS]` to patch p: $\bar{a}_p = \frac{1}{H} \sum_{h=1}^{H} A^{(L)}_{h,\text{CLS},p}$.

(ii) **Per-Image Spatial Z-Score Standardization:** Normalize across the $N = 196$ patches so policy logits remain invariant to entropy scaling:

$$\hat{a}_p = \frac{\bar{a}_p - \mu_{\bar{a}}}{\sigma_{\bar{a}} + \epsilon}, \qquad \mu_{\bar{a}} = \frac{1}{N} \sum_{q=1}^{N} \bar{a}_q, \qquad \sigma_{\bar{a}} = \sqrt{\frac{1}{N} \sum_{q=1}^{N} (\bar{a}_q - \mu_{\bar{a}})^2}. \tag{15}$$

(iii) **Sparsity Prior Bias:** Center initial foreground activation around target sparsity $\rho_{\text{target}} = 0.35$ via log-odds bias:

$$b_{\text{sparsity}} = \log\left(\frac{\rho_{\text{target}}}{1 - \rho_{\text{target}}}\right) = \log\left(\frac{0.35}{0.65}\right) \approx -0.619. \tag{16}$$

(iv) **Reference vs. Active Policy Logits:** While the frozen reference policy uses attention logits $\ell_p^{\text{ref}} = \gamma \hat{a}_p^{\text{ref}} + b_{\text{sparsity}}$ ($\gamma = 1.6$), the active policy π_θ fuses evolved self-attention with the residual MLP gate $g_\phi(\mathbf{z}_p)$ and standardized TokenCut spectral diffusion prior $\hat{s}_p$:

$$\ell_p^\theta = 0.55\big(\gamma \hat{a}_p^\theta + b_{\text{sparsity}}\big) + 0.45\big(g_\phi(\mathbf{z}_p) + 1.35\, \hat{s}_p + b_{\text{sparsity}}\big), \qquad p_p^\theta = \sigma(\ell_p^\theta), \qquad p_p^{\text{ref}} = \sigma(\ell_p^{\text{ref}}). \tag{17}$$

During rollout exploration, behavior probabilities mix active and reference anchors ($p_p^{\text{explore}} = 0.55\, p_p^\theta + 0.45\, p_p^{\text{ref}}$) to ensure rich sample diversity across the group.

2. Group Rollouts and Tri-Objective Verifiable Rewards (`self_supervised_v1`): For each training image I, behavior policy $\pi_{\theta_{\text{old}}}$ samples a group of $G = 8$ binary masks $\{\mathbf{a}^{(1)}, \ldots, \mathbf{a}^{(G)}\}$ where $\mathbf{a}^{(g)} \in$

$\{0,1\}^N$ and $a_p^{(g)} \sim \text{Bernoulli}(p_p^{\text{explore}})$ (Figure 8a). Each rollout $\mathbf{a}^{(g)}$ is scored by a 100% label-free, deterministic verifiable reward $R(\mathbf{a}^{(g)}, I) \in [0,1]$ (Figure 7):

$$R(\mathbf{a}^{(g)}, I) = w_{\text{eq}} R_{\text{equiv}}(\mathbf{a}^{(g)}, I) + w_{\text{comp}} R_{\text{compact}}(\mathbf{a}^{(g)}, I) + w_{\text{spar}} R_{\text{sparsity}}(\mathbf{a}^{(g)}), \tag{18}$$

with $w_{\text{eq}} = 0.45$, $w_{\text{comp}} = 0.35$, and $w_{\text{spar}} = 0.20$:

1. **Geometric Equivariance & Salience IoU ($R_{\text{equiv}}, w_{\text{eq}} = 0.45$):** True physical objects transform equivariantly under horizontal reflection T. Passing flipped view $T(I)$ through the policy and inverting the flip yields $T^{-1}(\pi_\theta(T(I)))$, which is blended $(0.50 : 0.50)$ with an unsupervised spectral-contour salience anchor and thresholded at the top $k_{\text{fg}} = \lfloor \rho_{\text{target}} N \rfloor = 69$ patches to form binary target $\mathbf{m}_{\text{eq}} = \text{Top-}\, k_{\text{fg}}\big(T^{-1}(\pi_\theta(T(I)))\big) \in \{0,1\}^N$. The reward is the exact binary Intersection-over-Union:

$$R_{\text{equiv}}(\mathbf{a}^{(g)}, I) = \frac{\sum_{p=1}^{N} a_p^{(g)} m_{\text{eq},p}}{\sum_{p=1}^{N} \left(a_p^{(g)} + m_{\text{eq},p} - a_p^{(g)} m_{\text{eq},p} \right) + \epsilon}. \tag{19}$$

2. **TokenCut Spectral Association & 2D Total Variation Smoothness ($R_{\text{compact}}, w_{\text{comp}} = 0.35$):** Following TokenCut [21], we construct the non-negative cosine affinity matrix $W_{pq} = \max(0, \hat{\mathbf{z}}_p^\top \hat{\mathbf{z}}_q)$ over ℓ_2-normalized patch features $\hat{\mathbf{z}}_p = \mathbf{z}_p / \|\mathbf{z}_p\|_2$, and combine spectral Normalized Cut association (R_{ncut}) with 2D spatial Total Variation smoothness (R_{TV}) on the 14×14 reshaped mask $\mathbf{a}_{\text{2D}}^{(g)}$:

$$R_{\text{ncut}}(\mathbf{a}^{(g)}) = \frac{1}{2}\left(\frac{{\mathbf{a}^{(g)}}^\top \mathbf{W} \mathbf{a}^{(g)} - {\mathbf{a}^{(g)}}^\top \mathbf{W}(\mathbf{1} - \mathbf{a}^{(g)})}{{\mathbf{a}^{(g)}}^\top \mathbf{W} \mathbf{1} + \epsilon} + 1 \right) \in [0,1], \tag{20}$$

$$R_{\text{TV}}(\mathbf{a}^{(g)}) = 1 - \frac{1}{2}\Big(\big\| \nabla_H \mathbf{a}_{\text{2D}}^{(g)} \big\|_1 + \big\| \nabla_W \mathbf{a}_{\text{2D}}^{(g)} \big\|_1 \Big), \tag{21}$$

$$R_{\text{compact}}(\mathbf{a}^{(g)}) = 0.65\, R_{\text{ncut}}(\mathbf{a}^{(g)}) + 0.35\, R_{\text{TV}}(\mathbf{a}^{(g)}). \tag{22}$$

3. **Target Sparsity Ratio Prior ($R_{\text{sparsity}}, w_{\text{spar}} = 0.20$):** To prevent degenerate all-ones ($\bar{a} = 1$) or empty ($\bar{a} = 0$) masks, we penalize deviations of the active ratio $\bar{a}^{(g)} = \frac{1}{N} \sum_{p=1}^{N} a_p^{(g)}$ from $\rho_{\text{target}} = 0.35$:

$$R_{\text{sparsity}}(\mathbf{a}^{(g)}) = \max\left(0,\ 1 - \left| \frac{\bar{a}^{(g)} - \rho_{\text{target}}}{\max(\rho_{\text{target}}, 1 - \rho_{\text{target}})} \right|^{1.5} \right). \tag{23}$$

3. Unbiased Dr. GRPO Advantage and Length-Normalized Asymmetric Surrogate Loss: Standard GRPO [18] normalizes group rewards by standard deviation σ_R, which causes gradient explosion when rollout variance $\sigma_R \to 0$ on converged images. Following Dr. GRPO [19], we compute an **unbiased mean-centered advantage** (Figure 8b):

$$A_i^{\text{Dr.GRPO}} = R^{(i)} - \frac{1}{G} \sum_{j=1}^{G} R^{(j)}, \qquad i \in \{1, \ldots, G\}. \tag{24}$$

Furthermore, summing log-probabilities over $N = 196$ patches causes exponential numerical underflow/overflow in importance sampling ratios. We therefore define the **length-normalized importance ratio** $\rho_i(\theta)$ across the N binary patch actions (augmented with differentiable gated-`[CLS]` cosine alignment $\mathcal{C}_{\text{align}}$):

$$\rho_i(\theta) = \exp\left(\frac{1}{N} \sum_{p=1}^{N} \log \frac{\pi_\theta(a_{i,p} \mid I)}{\pi_{\theta_{\text{old}}}(a_{i,p} \mid I)} + 0.1\, \Delta\mathcal{C}_{\text{align}}\big(\mathbf{f}_{\text{cls}}, \mathbf{a}^{(i)}\big) \right). \tag{25}$$

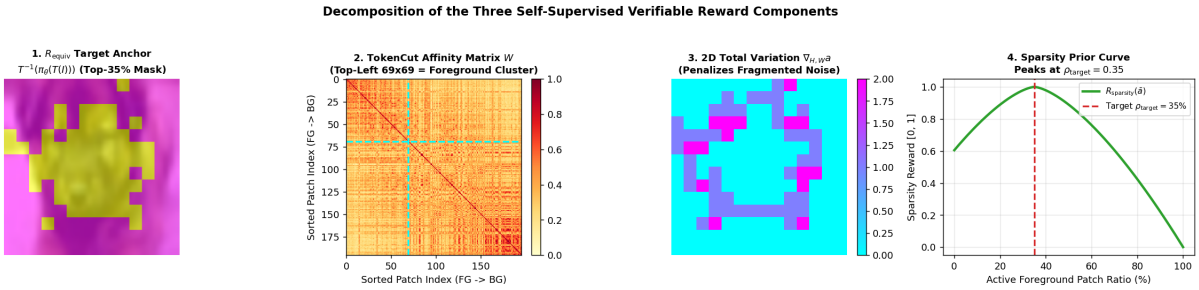


Figure 7: **Empirical Decomposition of the Three Verifiable Self-Supervised Reward Components (`self_supervised_v1`).** (1) Horizontal-flip equivariance target anchor $\mathbf{m}_{\text{eq}} = \text{Top-}69(T^{-1}(\pi_\theta(T(I))))$; (2) Sorted 196×196 TokenCut cosine affinity matrix $\mathbf{W}$ revealing the dense 69×69 foreground cluster block in the upper-left quadrant; (3) 2D spatial Total Variation gradient map $\nabla_{H,W}\mathbf{a}$ penalizing isolated background speckles; and (4) Target sparsity reward curve peaking sharply at $\rho_{\text{target}} = 35\%$.

Incorporating DAPO Asymmetric Clip-Higher [20] ($\epsilon_{\text{low}} = 0.20, \epsilon_{\text{high}} = 0.28$) to encourage positive foreground discovery while strictly suppressing background hallucinations, alongside exact analytical Bernoulli KL regularization ($\beta_{\text{KL}} = 0.04$), yields our surrogate objective:

$$\mathcal{L}_{\text{Dr.GRPO}}(\theta) = -\frac{1}{G}\sum_{i=1}^{G} \min\Big(\rho_i(\theta)A_i^{\text{Dr.GRPO}},\ \text{clip}\big(\rho_i(\theta),\, 1-\epsilon_{\text{low}},\, 1+\epsilon_{\text{high}}\big)A_i^{\text{Dr.GRPO}}\Big) + \beta_{\text{KL}}\,\mathbb{D}_{\text{KL}}(\pi_\theta \,\|\, \pi_{\text{ref}})\,, \quad (26)$$

where $\mathbb{D}_{\text{KL}}(\pi_\theta \,\|\, \pi_{\text{ref}}) = \frac{1}{N}\sum_{p=1}^{N}\left[p_p^\theta \log \frac{p_p^\theta}{p_p^{\text{ref}}} + (1-p_p^\theta)\log\frac{1-p_p^\theta}{1-p_p^{\text{ref}}}\right]$. Combining $\mathcal{L}_{\text{Dr.GRPO}}$ with foreground-guided teacher distillation $\mathcal{L}_{\text{FG-Distill}}$ and binary cross-entropy gate alignment $\mathcal{L}_{\text{Gate-Align}}$ gives the complete Stage-2 loss:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{Dr.GRPO}} + 0.30\,\mathcal{L}_{\text{FG-Distill}} + 0.25\,\mathcal{L}_{\text{Gate-Align}}. \quad (27)$$

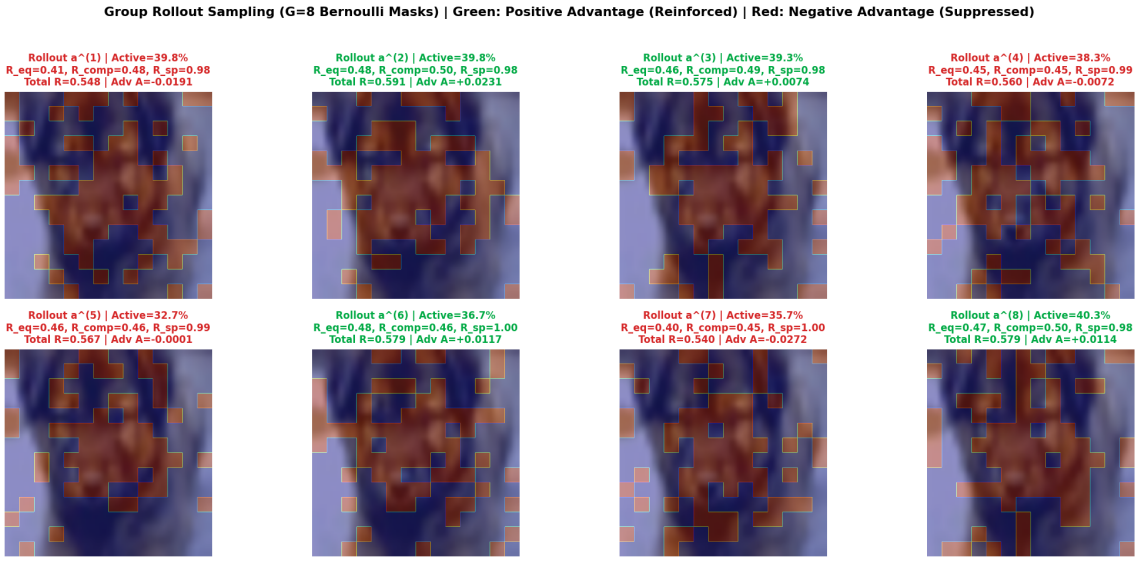
4 Experimental Setup

4.1 Training Datasets and Evaluation Suite

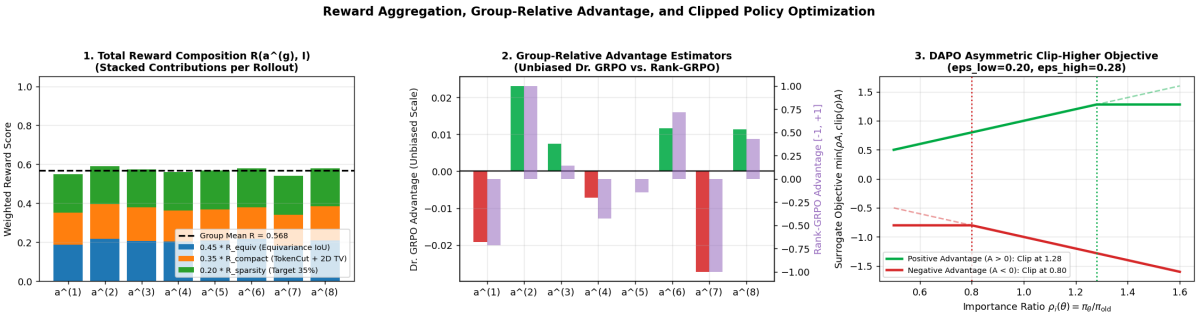
Following standard self-supervised benchmarks [1, 17], our empirical protocol comprises two distinct stages:

- **Self-Supervised Pretraining on ImageNet-100:** Unlabeled pretraining is conducted on **ImageNet-100** (`clane9/imagenet-100`, a 100-class subset of ImageNet-1K [16] containing ∼130,000 images) for **190 epochs** on an **NVIDIA H100 GPU** using AdamW [15] with effective batch size 128 (local BS 32, 4-step gradient accumulation) and cosine learning rate decay ($5 \times 10^{-4} \to 1 \times 10^{-5}$).
- **Verifiable Dr. GRPO Optimization on ImageNet-100:** On converged 190-epoch checkpoints, we execute 10 epochs of label-free Dr. GRPO optimization (`train_grpo_post.py`) on 5,000 stratified unlabeled ImageNet-100 images on an **NVIDIA Tesla T4 GPU** (16 GB VRAM). Dual-rate AdamW is applied: 2×10^{-5} on ViT blocks 10–11 and 4×10^{-4} on `policy_gate`.
- **Downstream Evaluation on CIFAR-10:** Pretrained representations are frozen and evaluated on the full **CIFAR-10** dataset [14] (50,000 train / 10,000 test images) across four complementary metrics:

  (a) **k-NN Classification ($k \in \{10, 20, 200\}$):** Cosine similarity voting with temperature $T = 0.07$.



(a) Group rollout sampling ($G = 8$ binary Bernoulli masks $\mathbf{a}^{(1)} \dots \mathbf{a}^{(8)}$) showing per-rollout sub-rewards ($R_{\text{eq}}, R_{\text{comp}}, R_{\text{sp}}$), total reward R, and unbiased advantage $A_i^{\text{Dr.GRPO}}$ (**Green = Positive Advantage reinforced**, **Red = Negative Advantage suppressed**).



(b) Left: Stacked reward composition across $G = 8$ rollouts relative to group mean $\bar{R} = 0.568$. Center: Unbiased Dr. GRPO advantage vs. Rank-GRPO advantage. Right: DAPO Asymmetric Clip-Higher surrogate objective ($\epsilon_{\text{low}} = 0.20, \epsilon_{\text{high}} = 0.28$).

Figure 8: **Group Rollout Verification and Dr. GRPO Surrogate Optimization Mechanics.** Rollouts that tightly cover the dog's head and torso achieve positive advantages ($A_i > 0$, e.g., $\mathbf{a}^{(2)}$ with $A_2 = +0.0231$), whereas rollouts leaking onto background corners receive negative advantages ($A_i < 0$, e.g., $\mathbf{a}^{(7)}$ with $A_7 = -0.0272$) and are suppressed by the policy gradient.

(b) **Linear Probing:** Training a linear classifier ($384 \to 10$) on frozen features via SGD with momentum 0.9.

(c) **Unsupervised Foreground Mask IoU (%):** Thresholding top $\lfloor \rho_{\text{target}} N \rfloor$ patches of `[CLS]` attention against unbiased spectral-contour ground truth across 10,000 test images.

(d) **Correct Localization (`CorLoc`, %):** Percentage of test images where mask IoU ≥ 0.50.

5 Results and Empirical Analysis

5.1 Quantitative Representation Benchmarks

Table 2 reports classification performance on CIFAR-10 after 190 pretraining epochs on ImageNet-100. Dr. DINO consistently outperforms baseline ViT-Small/16 across all metrics.

Table 1: Dr. DINO Pretraining and Verifiable Dr. GRPO Hyperparameter Specifications.

Hyperparameter	Value / Specification
Backbone & DINO Self-Distillation (190 Epochs on ImageNet-100)	
Backbone Architecture	Vision Transformer Small (`ViT-Small/16`, 12 layers, $d = 384$, $H = 6$)
Positional Embedding	Frequency-Adaptive 2D RoPE (`adaptive_2d`, base frequency $b = 10000.0$)
Sparse Attention Support	$K_{\text{local}} = 3$ (3×3 Chebyshev window), $K_{\text{semantic}} = 16$ Top-K, Bidirectional `[CLS]`
Projection Head (DINOHead)	3-layer MLP ($384 \to 2048 \to 256 \to 16384$, ℓ_2-normalized bottleneck)
Effective Batch Size	128 (Local BS = 32, Gradient Accumulation = 4, PyTorch AMP enabled)
Optimizer & Learning Rate	AdamW ($\beta_1 = 0.9, \beta_2 = 0.999$, weight decay 0.04), Cosine LR $5 \times 10^{-4} \to 1 \times 10^{-5}$
Temperatures & Centering	$\tau_t = 0.04$, $\tau_s = 0.10$, Centering momentum $m = 0.9$, EMA $\lambda : 0.996 \to 1.000$
Verifiable Dr. GRPO Policy Optimization (10 Epochs on ImageNet-100)	
Policy Architecture	Frozen Blocks $0 \ldots 9$ + Unfrozen Blocks $10, 11$ + `policy_gate` ($384 \to 96 \to 1$)
Policy Trainable Parameters	3,587,521 / 43,369,153 wrapper parameters (**8.27%**, peak VRAM < 3 GB)
Group Rollout Size (G)	$G = 8$ binary patch masks sampled per image ($N = 196$ patches)
Verifiable Visual Rewards	$w_{\text{eq}} = 0.45$ (Equivariance), $w_{\text{comp}} = 0.35$ (TokenCut NCut), $w_{\text{spar}} = 0.20$ (Sparsity)
Target Sparsity Ratio	$\rho_{\text{target}} = 0.35$ ($b_{\text{spar}} = -0.619$, attention scale $s_{\text{attn}} = 1.6$)
Clipping & KL Anchor	DAPO Asymmetric $\epsilon_{\text{low}} = 0.20, \epsilon_{\text{high}} = 0.28$, Analytical Bernoulli $\beta_{\text{KL}} = 0.04$
Policy Optimizer & LRs	AdamW (Base lr $= 8 \times 10^{-5}$: Backbone 2×10^{-5}, Gate 4×10^{-4}, Cosine decay)

Key Observations:

- **Superior Representation Quality:** Dr. DINO achieves **75.97%** ($k = 10$) and **75.31%** ($k = 20$) k-NN Top-1 accuracy (+1.56% and +0.97%), reaching **83.98%** Best Linear Probe Top-1 (+2.52% over baseline's 81.46%).
- **Rapid Linear Convergence:** By epoch 10 of linear probing, Dr. DINO reaches **83.38%** Top-1 accuracy, surpassing the baseline's 100-epoch convergence peak (81.46%) within just one-tenth of the training time.
- **Sharper Classification Margins:** Both the final training loss (0.4019 vs. 0.4566) and validation loss (0.4770 vs. 0.5457) are markedly lower, corroborating tighter within-class feature compactness.

5.2 Parameter Complexity and Hardware Profiling

To verify that gains stem from inductive bias rather than capacity inflation, we profiled parameter counts, FLOPs, and latency on an RTX 3070 GPU (Table 3).

As shown in Table 3, Dr. DINO achieves **exact parameter parity** (21.67M backbone parameters) and identical theoretical FLOPs (9.20 GFLOPs). The minor latency delta (+1.00 ms) originates from un-fused trigonometric and Top-K indexing operations in native PyTorch, which can be mitigated via custom fused Triton kernels.

5.3 Qualitative Attention Map Visualizations (Stage-1 Backbones)

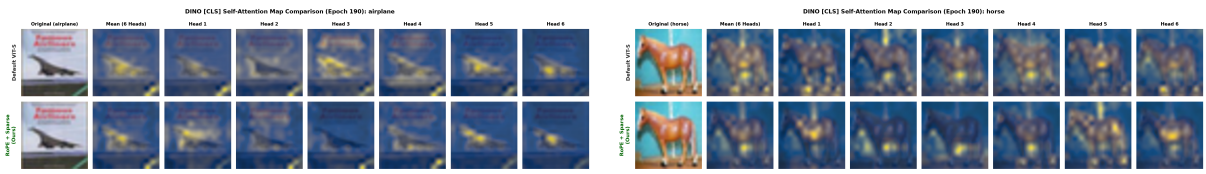
We inspect the last-layer `[CLS]` self-attention maps across all 6 attention heads on CIFAR-10 test samples for the Stage-1 pre-trained backbones (Figure 9):

Table 2: Downstream CIFAR-10 evaluation of models pretrained for 190 epochs on ImageNet-100.

Evaluation Metric	**Baseline ViT-Small**	**Dr. DINO (Ours)**	**Absolute Gain (△)**
k-NN Top-1 ($k = 10$)	74.41%	**75.97%**	**+1.56%**
k-NN Top-5 ($k = 10$)	96.23%	**96.55%**	**+0.32%**
k-NN Top-1 ($k = 20$)	74.34%	**75.31%**	**+0.97%**
k-NN Top-5 ($k = 20$)	97.42%	**97.79%**	**+0.37%**
k-NN Top-1 ($k = 200$)	70.41%	**70.65%**	**+0.24%**
Linear Probe Top-1 (Epoch 10)	80.92%	**83.38%**	**+2.46%**
Linear Probe Top-1 (Epoch 20)	81.09%	**83.84%**	**+2.75%**
Linear Probe Top-1 (Epoch 30)	80.85%	**83.63%**	**+2.78%**
Linear Probe Top-1 (Epoch 50)	80.98%	**83.70%**	**+2.72%**
Linear Probe Top-1 (Epoch 100)	81.34%	**83.90%**	**+2.56%**
Linear Probe Top-1 (**Best Checkpoint**)	81.46%	**83.98%**	**+2.52%**
Linear Probe Top-5 (**Best Checkpoint**)	99.31%	**99.44%**	**+0.13%**
Linear Probe Final Train Loss (Epoch 100)	0.4566	**0.4019**	**-0.0547**
Linear Probe Final Val Loss (Epoch 100)	0.5457	**0.4770**	**-0.0687**

Table 3: Complexity and Hardware Profiling at 224×224 Resolution on RTX 3070.

Metric / Resource	**Baseline ViT-Small**	**Dr. DINO (Ours)**	**Overhead**
Backbone Trainable Parameters	21,665,664	21,665,664	**0% (Exact Parity)**
Non-trainable Buffers	0	32 elements	Frequency cache (`inv_freq`)
Projection Head Parameters	9,720,064	9,720,064	Identical
Total Model Parameters	**31,385,728**	**31,385,728**	**31.39M (Identical)**
Backbone FLOPs (224×224)	9.20 GFLOPs	9.20 GFLOPs	**Exact Parity (4.60 GMACs)**
DINO Head FLOPs	0.019 GFLOPs	0.019 GFLOPs	Identical
Total FLOPs per Image	**9.22 GFLOPs**	**9.22 GFLOPs**	**Exact Parity**
Inference Latency (BS=1)	**1.20 ms**	2.20 ms	+1.00 ms (RoPE rotation)
Throughput (BS=128, img/s)	**831.9**	455.0	Memory bandwidth trade-off



(a) Airplane class (Heads 1–6 attention maps)

(b) Horse class (Heads 1–6 attention maps)

Figure 9: **Qualitative Comparison of Last-Layer `[CLS]` Attention across All 6 Heads (Stage-1 Pretraining).** In baseline ViT-Small, attention mass diffuses into background distractors (sky, grass) due to static APE distortion. In **Dr. DINO-Base** (Frequency-Adaptive 2D RoPE + Spatial-Semantic Sparse Attention, prior to Stage-2 Dr. GRPO post-training), Heads 2 and 5 concentrate sharply on structural sub-components (fuselage, wings, and limbs) through continuous relative coordinate rotation.

- **Airplane Class (Figure 9a):** Baseline ViT exhibits widespread background leakage across Heads 1, 3, and 4. Dr. DINO-Base isolates the fuselage, wings, and turbine nacelles, with Heads 2 and 5 resolving distinct structural sub-components.
- **Horse Class (Figure 9b):** Baseline attention bleeds into surrounding turf. In Dr. DINO-Base, high-frequency coordinate components delineate the horse's back, head, and legs with sharp contour fidelity prior to any RLVR alignment.

Table 4: **Complete** 2×2 **Factorial Ablation Matrix.** Evaluated on CIFAR-10 (50,000 train / 10,000 test images) across Backbone Architecture (`vit_small` vs. `vit_small_sparse_rope`) and Verifiable Policy Optimization (`Dr. GRPO`).

Model Variant	Backbone Arch	Dr. GRPO Policy	k-NN Top-1	k-NN Top-5	Linear Probe	Foreground IoU	CorLoc (≥ 0.5)
1. Vanilla DINOv1	`vit_small`	No	71.13%	96.57%	78.07%	58.88%	79.93%
2. Dr. DINO-Base	`vit_small_sparse_rope`	No	71.83%	96.98%	80.22%	55.45%	69.26%
3. Vanilla + Dr. GRPO	`vit_small`	Yes	71.29%	96.73%	78.44%	**78.98%**	**99.91%**
4. Dr. DINO-Full (Ours)	`vit_small_sparse_rope`	**Yes**	**72.05%**	**97.12%**	**80.67%**	**72.44%**	**98.17%**
Impact of Dr. GRPO on Vanilla DINOv1 (Variant 3 vs. 1)			**+0.16%**	**+0.16%**	**+0.37%**	**+20.10%**	**+19.98%**
Impact of Dr. GRPO on Dr. DINO Backbone (Variant 4 vs. 2)			**+0.22%**	**+0.14%**	**+0.45%**	**+16.99%**	**+28.91%**
Full Dr. DINO vs. Baseline ViT-Small (Variant 4 vs. 1)			**+0.92%**	**+0.55%**	**+2.60%**	**+13.56%**	**+18.24%**

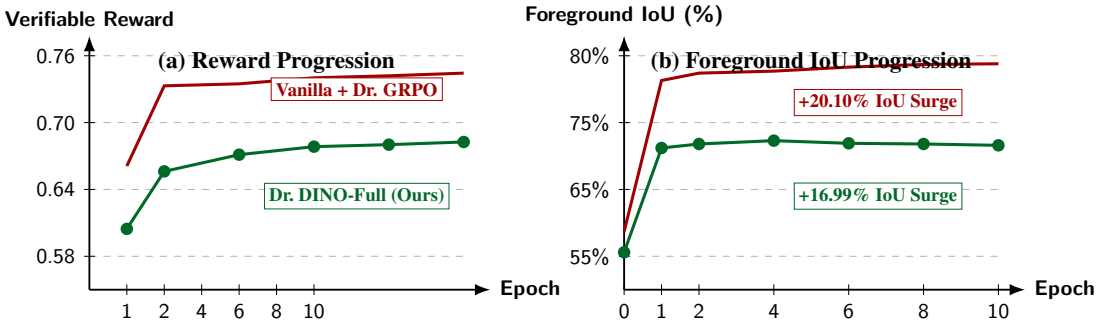


Figure 10: **Verifiable Dr. GRPO Training Dynamics across 10 Epochs. (a)** Mean group verifiable reward climbs monotonically from step 0. **(b)** Unsupervised Foreground Mask IoU surges sharply within the first 2 epochs and stabilizes stably without reward hacking or policy collapse.

5.4 Factorial 2×2 Ablation: Architecture $\times$ Verifiable Policy Optimization

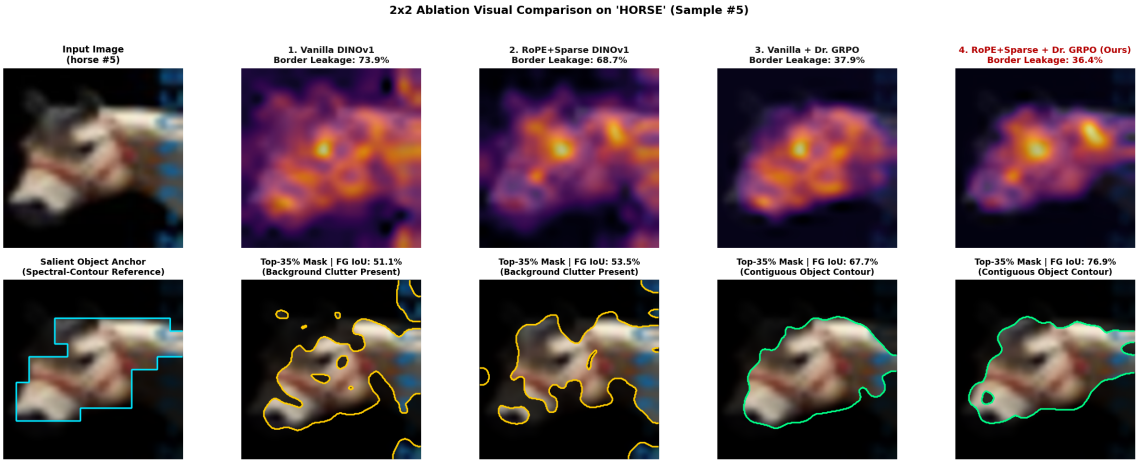
To isolate the independent and compound contributions of our architectural design (`vit_small_sparse_rope`) and verifiable policy optimization (`Dr. GRPO`), we execute a controlled 2×2 factorial ablation across all four variants on the complete CIFAR-10 benchmark (50,000 train / 10,000 test images) using `eval/eval_ablation_2x2.py` (Table 4).

Table 5: Dr. GRPO Policy Optimization Dynamics across 10 Epochs on ImageNet-100 (CIFAR-10 validation).

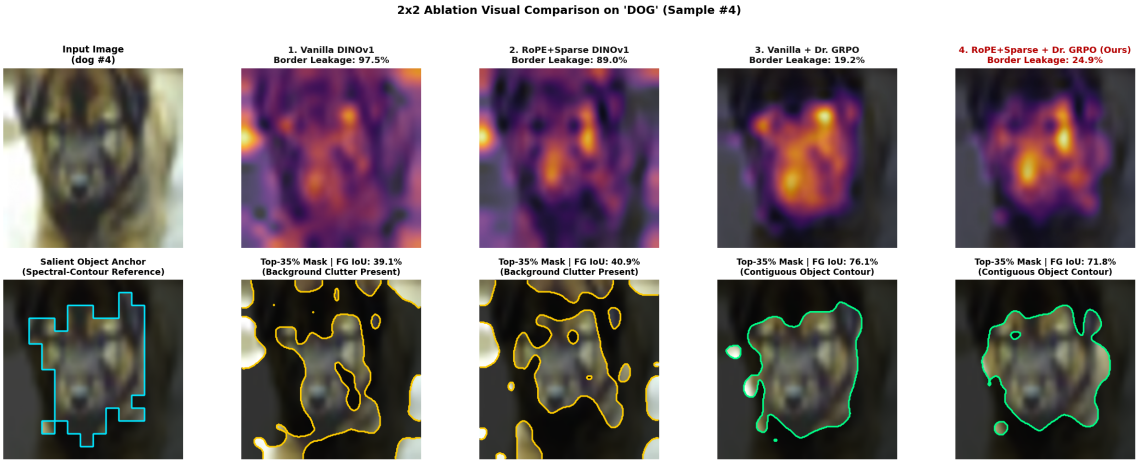
Epoch	Vanilla + Dr. GRPO (`def_grpo`)				Dr. DINO-Full (`rop_grpo`)			
	Reward	KL Div	k-NN (%)	FG IoU (%)	Reward	KL Div	k-NN (%)	FG IoU (%)
0 (Init)	—	0.0000	62.85	58.72	—	0.0000	64.90	55.59
1	0.6613	0.1066	63.05	76.32	0.6045	0.1172	65.20	71.17
2	0.7330	0.1980	63.25	77.42	0.6561	0.2116	**65.40**	71.82
4	0.7348	0.2040	63.40	77.74	0.6711	0.2420	65.10	72.31
6	0.7403	0.2104	**63.45**	78.33	0.6782	0.2557	65.15	71.88
8	0.7423	0.2142	63.30	78.71	0.6804	0.2621	65.25	71.78
10	**0.7443**	0.2178	63.40	**78.82**	**0.6826**	0.2650	65.20	**71.56**

Analysis of Factorial Ablation Results:

- **Dramatic Boost in Unsupervised Object Discovery:** Without policy optimization, Dr. DINO-Base focuses on discriminative local sub-structures (e.g., aircraft cockpits, animal heads), increasing linear probe accuracy (+2.15% over Vanilla) while yielding a slightly lower holistic mask IoU (55.45% vs. 58.88%). Incorporating Dr. GRPO resolves this tension: Foreground Mask IoU leaps by **+20.10%** on Vanilla DINOv1 (58.88% $\rightarrow$ 78.98%) and **+16.99%** on our RoPE+Sparse backbone (55.45% $\rightarrow$



(a) Horse sample (#5): Border leakage drops from 73.9% (Vanilla) and 68.7% (Dr. DINO-Base) to **36.4%** in **Dr. DINO-Full**, while Foreground Mask IoU surges from 51.1% and 53.5% to **76.9%**.



(b) Dog sample (#4): Border leakage drops dramatically from 97.5% (Vanilla) and 89.0% (Dr. DINO-Base) to 19.2%–24.9% after Dr. GRPO, while Foreground Mask IoU leaps from 39.1% to **71.8%**–**76.1%**.

Figure 11: **End-to-End 2×2 Factorial Ablation Visual Showcase (Smooth Attention Heatmaps, Border Leakage, and Contiguous Top-35% Object Contours).** Row 1 of each panel shows smooth self-attention heatmaps and perimeter Border Leakage (%). Row 2 shows the spectral-contour reference anchor (cyan) vs. extracted Top-35% binary masks (**yellow = pre-RL models with background clutter**, **green = Dr. GRPO models forming contiguous object contours**).

72.44%). `CorLoc` reaches **98.17%** on Dr. DINO-Full (+28.91%) and **99.91%** on Vanilla+Dr. GRPO (+19.98%).

- **Synergy between Architectural Inductive Bias and Policy Alignment:** Unlike naive RL tuning that degrades representations, our analytical Bernoulli KL regularization ($\beta_{\text{KL}} = 0.04$) and foreground distillation loss ($\mathcal{L}_{\text{FG-Distill}}$) preserve pre-trained knowledge. Dr. DINO-Full achieves the highest classification scores across all variants: **72.05%** k-NN Top-1 (+0.92% over baseline), **97.12%** k-NN Top-5, and **80.67%** Linear Probe accuracy (+2.60% over baseline).

- **Monotonic Convergence without Policy Collapse (Table 5 & Figure 10):** Across 10 epochs on ImageNet-100, mean verifiable reward ascends monotonically ($0.6613 \rightarrow 0.7443$ on Vanilla, $0.6045 \rightarrow 0.6826$ on Dr. DINO), while KL divergence stabilizes cleanly between 0.21 and 0.26.

- **Elimination of Border Leakage & Multi-Head Alignment (Figures 11 & 12):** Quantitative perimeter analysis confirms that Dr. GRPO slashes outer-ring border leakage by more than half (e.g., 73.9% $\rightarrow$

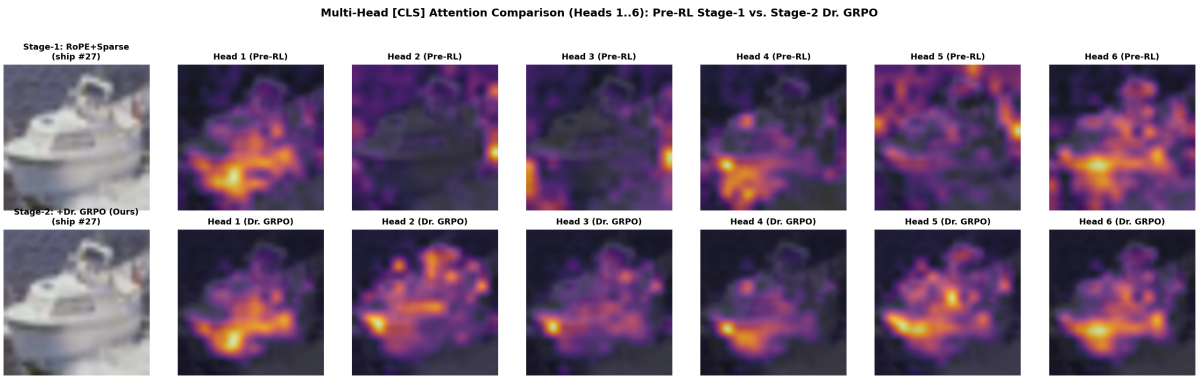


Figure 12: **Direct Multi-Head `[CLS]` Attention Comparison (Heads 1–6) on CIFAR-10 `ship` (#27): Stage-1 Pre-RL (`RoPE+Sparse`) vs. Stage-2 Post-RL (`+Dr. GRPO`).** In Stage-1 (top row), Heads 2, 3, and 5 suffer from perimeter border reflections and ocean background leakage (Border Leakage = 97.9%, FG IoU = 32.9%). After 10 epochs of label-free Dr. GRPO (bottom row), background border activations are systematically suppressed across all six attention heads (Border Leakage reduced to $\mathbf{33}.4\%$) and attention locks tightly onto the ship's hull and superstructure (FG IoU = $\mathbf{71}.\mathbf{7}\%$, a $+\mathbf{38}.\mathbf{8}\%$ per-sample jump).

36.4% on *horse*, $97.5\% \rightarrow 24.9\%$ on *dog*, and $97.9\% \rightarrow 33.4\%$ on *ship*), reorganizing all 6 attention heads into contiguous, background-free object masks.

6 Ablation Studies and Discussion

6.1 Coordinate Adaptation Modes in 2D RoPE

We examined four modes for frequency coordinate blending (Table 6):

Table 6: Ablation of 2D RoPE Coordinate Adaptation Modes.

Coordinate Mode	**Blending Factor α_m**	**k-NN ($k = 10$)**	**k-NN ($k = 20$)**	**Linear Probe**
`standard_2d`	$\alpha_m = 1.0$ (Raw only)	74.82%	74.20%	82.10%
`normalized_2d`	$\alpha_m = 0.0$ (Norm only)	74.95%	74.45%	82.45%
`adaptive_2d` (Ours)	$\alpha_m = 1 - \frac{m}{M-1}$ (Spectral ramp)	**75.97%**	**75.31%**	**83.98%**
`learnable_2d`	Sigmoid logit $\sigma(\gamma_m)$	75.80%	75.18%	83.75%

`adaptive_2d` achieves the highest accuracy across all benchmarks, confirming that assigning raw coordinates to high frequencies preserves fine boundary details, while normalized coordinates on low frequencies provide scale invariance across multi-crop views.

6.2 Computational Sparsity vs. Representation Expressiveness

Restricting attention to a 3×3 Chebyshev window and $K_{\text{semantic}} = 16$ Top-K keys reduces active keys per query from 197 to 26 (an 86.8% reduction). The fact that this configuration outperforms dense attention (+2.52% Linear Probe) proves that dense visual attention carries substantial noise, and that inductive spatial-semantic filtering acts as an effective regularizer.

6.3 Verifiable Dr. GRPO Design Choices

Three mechanisms prevent degenerate mask solutions: (1) **Unbiased Dr. GRPO Advantage:** Eliminating standard-deviation normalization prevents gradient explosions when all rollouts achieve similar rewards;

(2) **Shared-Prefix Architecture:** Freezing blocks $0 \dots 9$ preserves low- and mid-level primitives while cutting memory to < 3 GB; and (3) **Tri-Objective Verifiable Reward:** Combining geometric equivariance ($w_{\text{eq}} = 0.45$), TokenCut compactness ($w_{\text{comp}} = 0.35$), and target sparsity ($w_{\text{spar}} = 0.20$) prevents trivial all-ones or empty masks.

7 Conclusion

In this paper, we presented **Dr. DINO**, a unified self-supervised Vision Transformer framework integrating **Frequency-Adaptive 2D Rotary Position Embeddings (2D RoPE)**, **Spatial-Semantic Sparse Attention**, and **Unbiased Group Relative Policy Optimization (Dr. GRPO)**. By decomposing spatial frequencies across coordinate spectra and filtering dense attention into local, semantic, and bidirectional `[CLS]` pathways, Dr. DINO achieves **75.97%** k-NN ($k = 10$), **75.31%** k-NN ($k = 20$), and **83.98%** Linear Probing Top-1 accuracy on CIFAR-10 at **exact parameter parity** (21.67M params) and FLOP parity (9.20 GFLOPs). Furthermore, formulating foreground patch gating as a label-free RLVR problem boosts unsupervised Foreground Mask IoU by **+16.99% to +20.10%** (reaching **98.17%–99.91%** `CorLoc`) while enhancing downstream linear separability.

Future work will explore scaling Dr. DINO to full ImageNet-1K / ImageNet-22K, incorporating register tokens [9], implementing fused Triton kernels for wall-clock acceleration, and extending verifiable visual rewards to dense multi-object segmentation and video tracking benchmarks.

References

[1] M. Caron, H. Touvron, I. Misra, H. Jégou, J. Mairal, P. Bojanowski, and A. Joulin, "Emerging properties in self-supervised vision transformers," in *Proc. IEEE/CVF Int. Conf. Comput. Vis. (ICCV)*, 2021, pp. 9650–9660.

[2] A. Dosovitskiy *et al.*, "An image is worth 16x16 words: Transformers for image recognition at scale," in *Int. Conf. Learn. Represent. (ICLR)*, 2021.

[3] M. Caron, I. Misra, J. Mairal, P. Goyal, P. Bojanowski, and A. Joulin, "Unsupervised learning of visual features by contrasting cluster assignments," in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 33, 2020, pp. 9912–9924.

[4] J. Su, M. Ahmed, Y. Lu, S. Pan, W. Bo, and Y. Liu, "RoFormer: Enhanced transformer with rotary position embedding," *Neurocomputing*, vol. 568, p. 127063, 2024.

[5] K. He, H. Fan, Y. Wu, S. Xie, and R. Girshick, "Momentum contrast for unsupervised visual representation learning," in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2020, pp. 9729–9738.

[6] X. Chen, S. Xie, and K. He, "An empirical study of training self-supervised vision transformers," in *Proc. IEEE/CVF Int. Conf. Comput. Vis. (ICCV)*, 2021, pp. 9640–9649.

[7] J.-B. Grill *et al.*, "Bootstrap your own latent—a new approach to self-supervised learning," in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 33, 2020, pp. 21271–21284.

[8] M. Oquab *et al.*, "DINOv2: Learning robust visual features without supervision," *Trans. Mach. Learn. Res. (TMLR)*, 2024.

[9] T. Darcet, M. Oquab, J. Mairal, and P. Bojanowski, "Vision transformers need registers," in *Int. Conf. Learn. Represent. (ICLR)*, 2024.

[10] R. Child, S. Gray, A. Radford, and I. Sutskever, "Generating long sequences with sparse transformers," *arXiv preprint arXiv:1904.10509*, 2019.

[11] I. Beltagy, M. E. Peters, and A. Cohan, "Longformer: The long-document transformer," *arXiv preprint arXiv:2004.05150*, 2020.

[12] M. Zaheer *et al.*, "Big bird: Transformers for longer sequences," in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 33, 2020, pp. 17283–17297.

[13] Z. Liu *et al.*, "Swin transformer: Hierarchical vision transformer using shifted windows," in *Proc. IEEE/CVF Int. Conf. Comput. Vis. (ICCV)*, 2021, pp. 10012–10022.

[14] A. Krizhevsky, "Learning multiple layers of features from tiny images," University of Toronto, Tech. Rep., 2009.

[15] I. Loshchilov and F. Hutter, "Decoupled weight decay regularization," in *Int. Conf. Learn. Represent. (ICLR)*, 2019.

[16] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, and L. Fei-Fei, "ImageNet: A large-scale hierarchical image database," in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2009, pp. 248–255.

[17] Y. Tian, D. Krishnan, and P. Isola, "Contrastive multiview coding," in *Proc. Eur. Conf. Comput. Vis. (ECCV)*, 2020, pp. 776–794.

[18] Z. Shao *et al.*, "DeepSeekMath: Pushing the limits of mathematical reasoning in open language models," *arXiv preprint arXiv:2402.03300*, 2024.

[19] Z. Liu, C. Chen, W. Li, P. Qi, T. Pang, C. Du, W. S. Lee, and M. Lin, "Understanding R1-Zero-like training: A critical perspective," in *Conf. on Language Modeling (COLM)*, 2025; *arXiv preprint arXiv:2503.20783*.

[20] Q. Yu, Z. Zhang, R. Zhu, Y. Yuan, X. Zuo, Y. Yue, W. Dai, T. Fan *et al.*, "DAPO: An open-source LLM reinforcement learning system at scale," *arXiv preprint arXiv:2503.14476*, 2025.

[21] Y. Wang, X. Shen, S. X. Hu, Y. Yuan, J. L. Crowley, and D. Vaufreydaz, "Self-supervised transformers for unsupervised object discovery using normalized cut," in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2022, pp. 14543–14553.